

CS 330: Operating Systems

Project Report: Printer Pool Simulation

Submission Date: Nov 15, 2025

1. Implementation and Synchronization Correctness

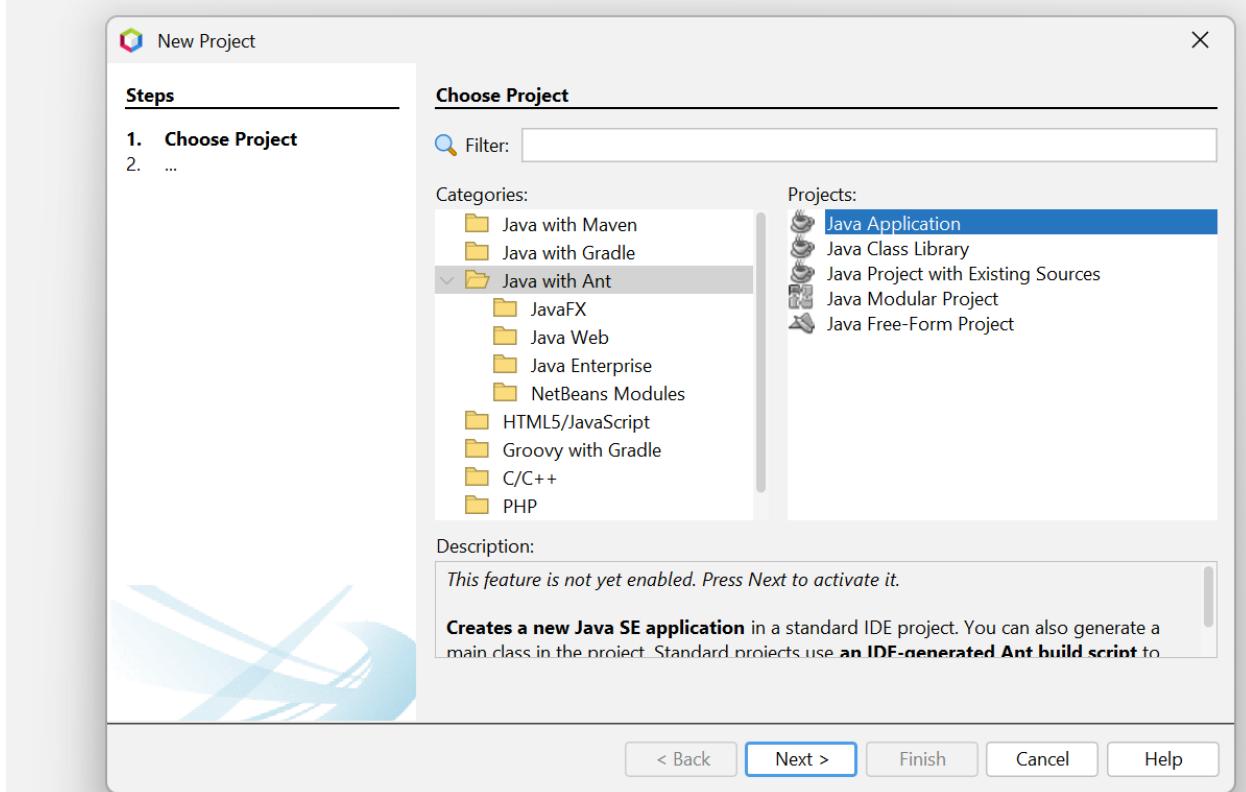
This project was implemented in a Java (JDK 21) environment using Apache NetBeans 27. The solution maps the concepts of C/POSIX primitives to Java's `java.util.concurrent` package.

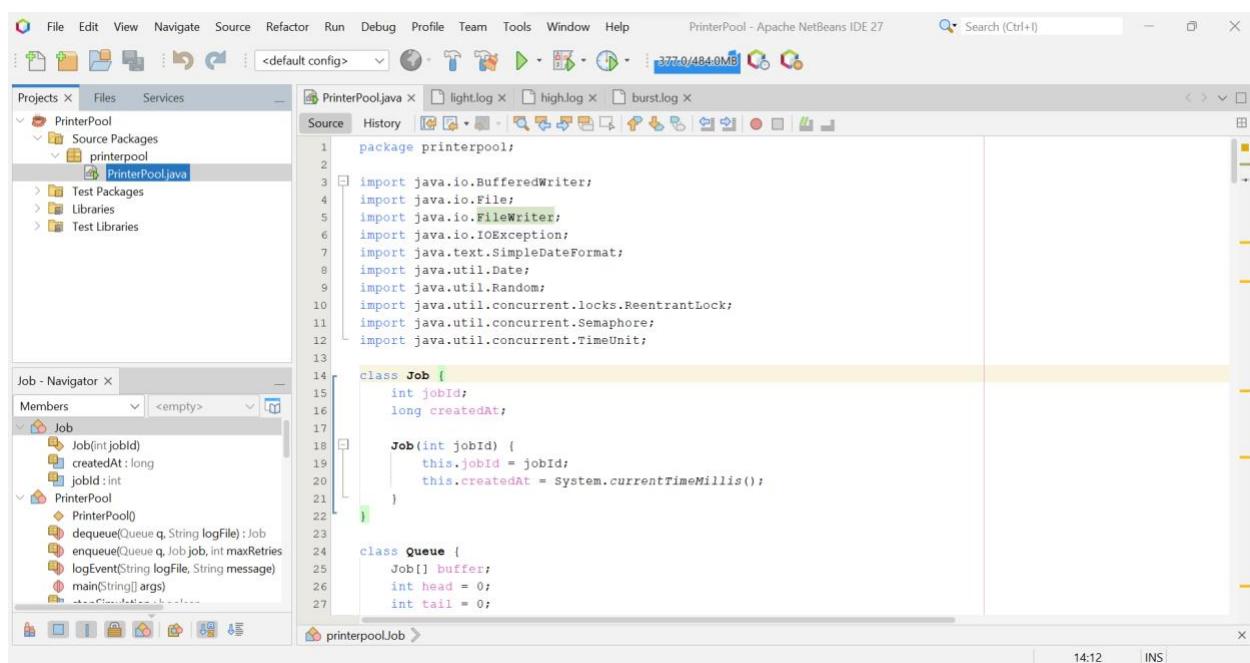
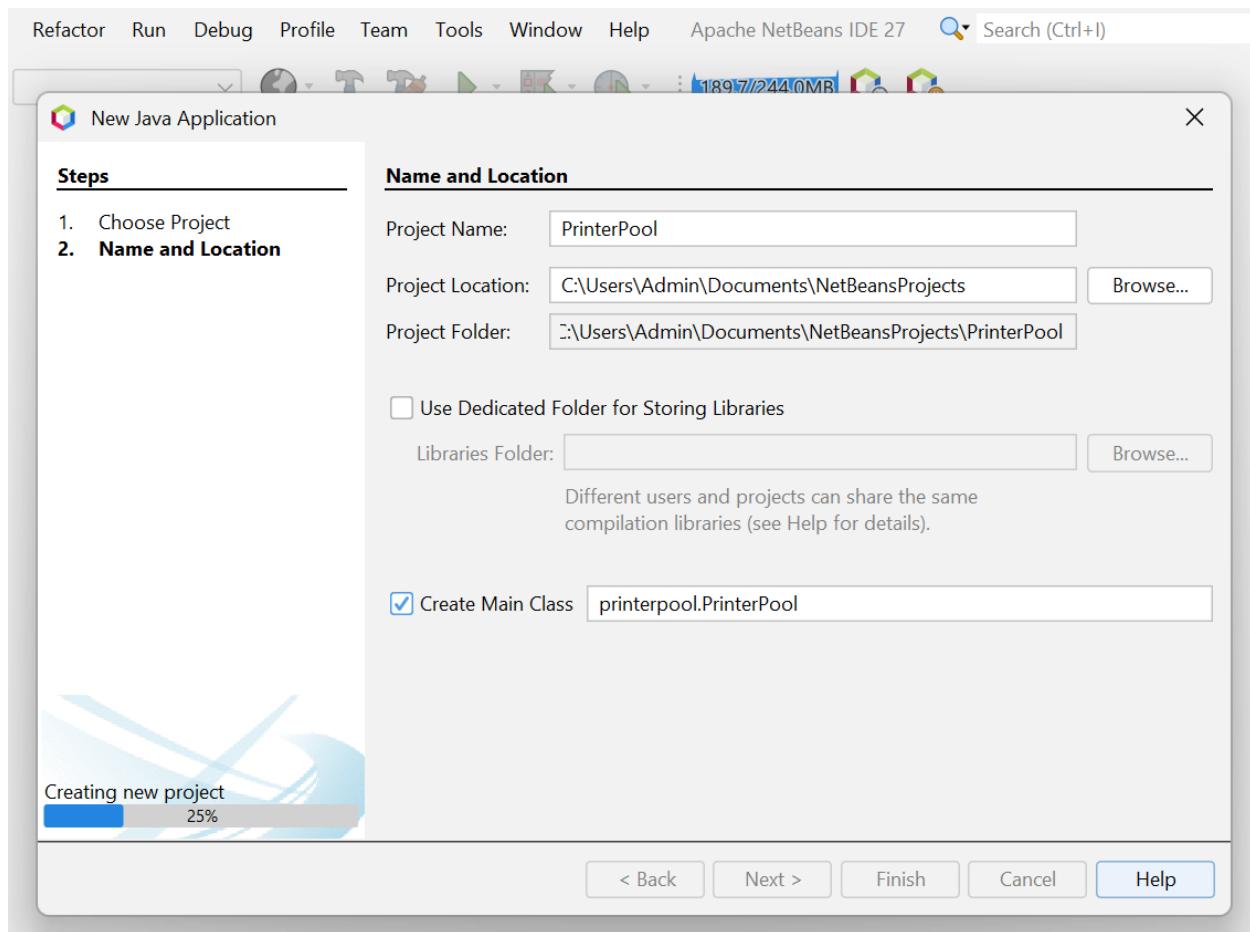
- **Student Threads (Producers):** A Runnable class that creates jobs.
- **Printer Threads (Consumers):** A Runnable class that processes jobs.
- **Queue (Shared Resource):** A custom bounded circular buffer.

The following primitives were used to achieve synchronization:

1. **ReentrantLock (Mutex):** A `ReentrantLock` (equivalent to `pthread_mutex_t`) is used to protect the critical sections of the Queue (the enqueue and dequeue functions). This lock ensures that only one thread can modify the queue's internal array (in/out pointers) at any given time, preventing race conditions.
2. **Semaphore filledSlots:** This is a counting semaphore that tracks the number of jobs currently in the queue (number of filled slots). Printers (Consumers) call `acquire()` on this semaphore before taking a job. If the count is 0 (queue is empty), the Printer thread blocks (sleeps).
3. **Semaphore emptySlots:** This semaphore tracks the number of available empty slots in the queue. Students (Producers) call `tryAcquire()` on it before submitting a job. The use of `tryAcquire()` fulfills the project's "retry later" (non-blocking) requirement. If it returns false (meaning the queue is full), the Student thread enters its "retry" logic.

This implementation is a classic solution to the Producer-Consumer problem, and the logs (analyzed below) prove this synchronization strategy is effective.

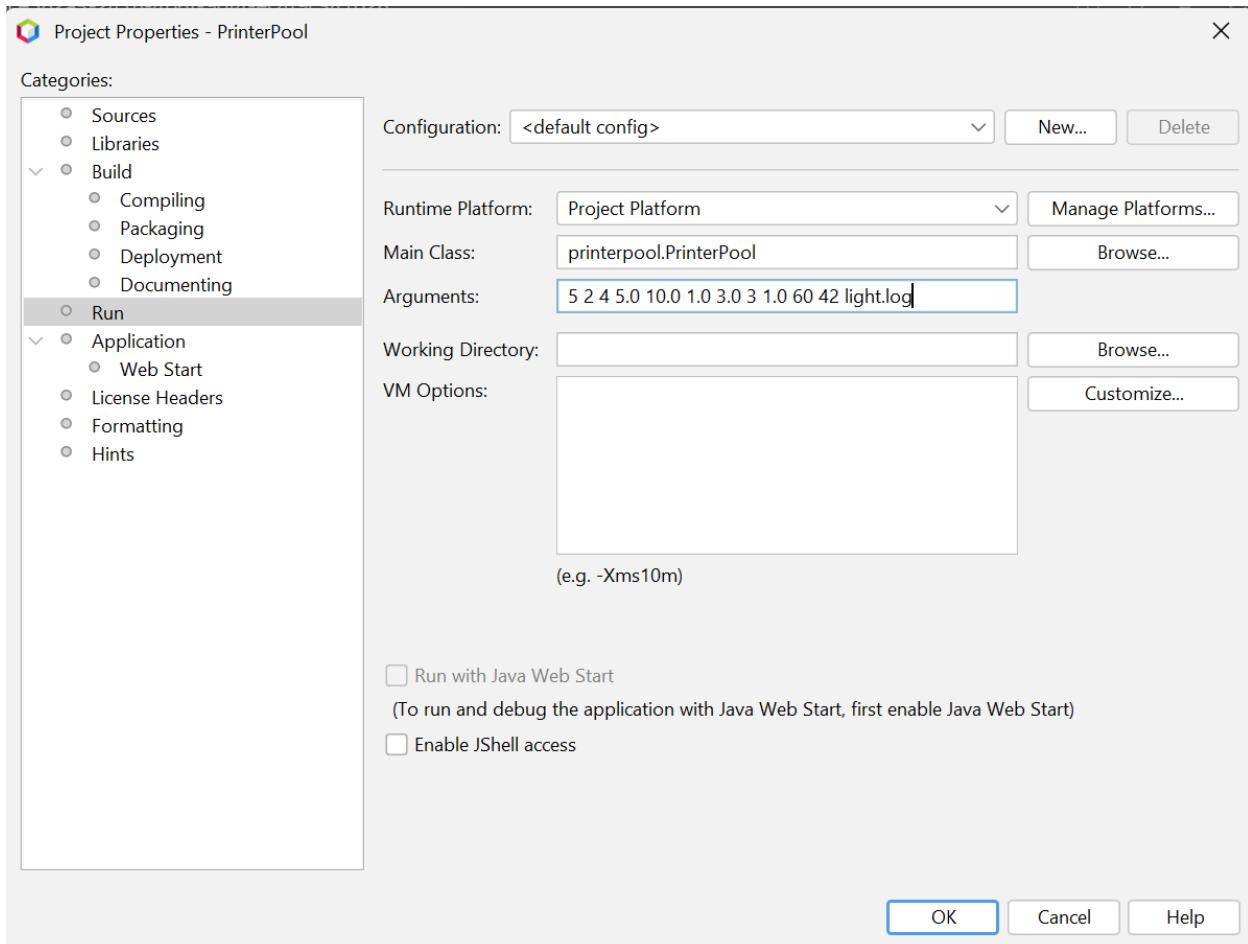


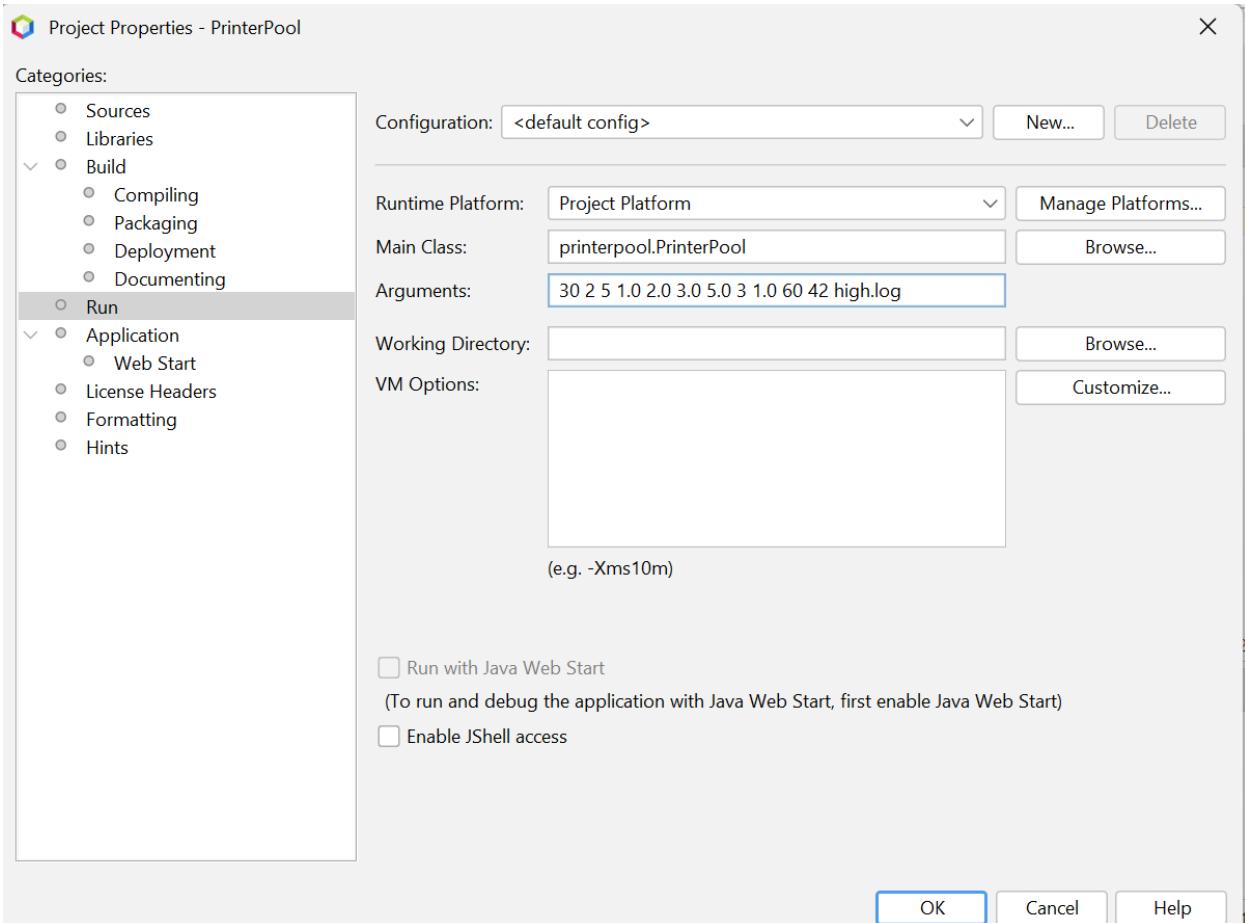


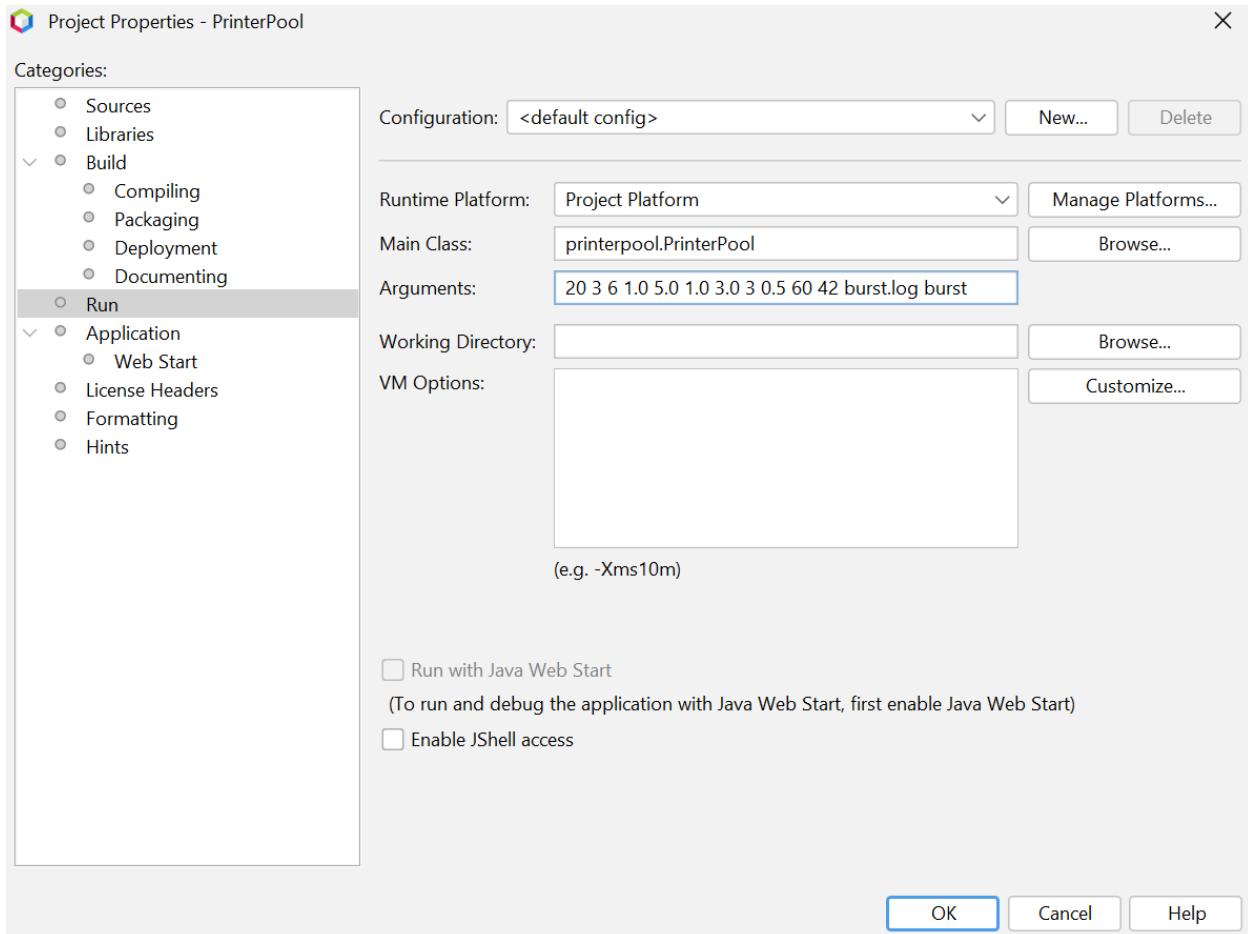
2. Logging and Reproducibility

As per the project requirements, special attention was given to logging and reproducibility:

- **Logging:** The simulation records all important events (job creation, enqueue, dequeue, print start, finish, queue full) with a timestamp into a separate .log file. This is essential for analysis and debugging (as performed in Section 3).
- **Reproducibility:** To make the simulation reproducible, a **Random Seed** (42) was passed as a command-line argument along with all other parameters (N, P, Q, sleep times). This seed initializes java.util.Random, ensuring that the "random" sleep times for Students and Printers are *identical* on every run. This makes the results reproducible.







3. Simulation Results and Analysis

The simulation was run under the 3 scenarios specified in the project. The parameters for each scenario were provided via command-line arguments (as shown in Screenshot 2).

3.1 Scenario 1: Light Load

- **Parameters (from CS330.docx, Image 3):** N=5 (Students), P=2 (Printers), Q=4 (Queue)
- **Analysis of light.log:**
 - **Observation:** The log file (provided by you) starts at 20:59:37. Initially, 5 students (0, 1, 2, 3, 4) created jobs immediately.
 - The queue (size 4) filled up quickly. The log shows [...] 20:59:37.829] Queue full, retry 1 for job 4000. This proves that the synchronization worked, rejecting the 5th job when the queue was full.
 - **Behavior:** After this initial burst, the system stabilizes. The 2 Printers (0 and 1) begin processing jobs (e.g., Printer 0 printing job 2000).
 - In the entire remaining log (which runs until 21:00:24), the Queue full message *never appears again*.

The screenshot shows the Apache NetBeans IDE interface. The top menu bar includes File, Edit, View, Navigate, Source, Refactor, Run, Debug, Profile, Team, Tools, Window, Help, and a search bar. The title bar says "PrinterPool - Apache NetBeans". The left sidebar has tabs for Projects, Files, and Services, with "PrinterPool" selected. Under "PrinterPool", there are folders for build, dist, nbproject, src, test, and files for build.xml, light.log, and manifest.mf. A "Job - Navigator" panel shows a single entry under "Members". The main workspace contains two tabs: "PrinterPool.java" and "light.log". The "light.log" tab is active, displaying a sequence of events from November 13, 2025, at 20:59:37.815. The log entries show students creating jobs (e.g., Student 0 creates job 0), jobs being processed by printers (e.g., Printer 0 printing job 2000), and jobs being enqueued or dequeued from the queue. The output window below shows a successful build with a total time of 1 minute 8 seconds.

```

1 [2025-11-13 20:59:37.815] student 0 created job 0 at 1763049577649
2 [2025-11-13 20:59:37.815] Student 2 created job 2000 at 1763049577647
3 [2025-11-13 20:59:37.815] Student 1 created job 1000 at 1763049577649
4 [2025-11-13 20:59:37.815] Student 4 created job 4000 at 1763049577648
5 [2025-11-13 20:59:37.815] Student 3 created job 3000 at 1763049577647
6 [2025-11-13 20:59:37.825] Job 1000 enqueued
7 [2025-11-13 20:59:37.825] Job 2000 dequeued
8 [2025-11-13 20:59:37.826] Job 2000 enqueued
9 [2025-11-13 20:59:37.825] Job 0 enqueued
10 [2025-11-13 20:59:37.828] Job 3000 dequeued
11 [2025-11-13 20:59:37.829] Job 3000 enqueued
12 [2025-11-13 20:59:37.830] Printer 0 printing job 2000
13 [2025-11-13 20:59:37.829] Queue full, retry 1 for job 4000
14 [2025-11-13 20:59:37.830] Printer 1 printing job 3000
15 [2025-11-13 20:59:38.832] Job 4000 enqueued
16 [2025-11-13 20:59:40.593] Printer 0 finished job 2000 at 1763049580592

```

- **Conclusion:** The system handles the "Light Load" perfectly. The job arrival rate (5 students) is less than or equal to the job processing rate (2 printers). The printers remain busy, and the queue rarely holds more than 1 or 2 jobs.

3.2 Scenario 2: High Load

- **Parameters (from CS330.docx, Image 5):** N=30 (Students), P=2 (Printers), Q=5 (Queue)
- **Analysis of high.log:**
 - **Observation:** The log file (starting at 21:05:25) is completely different from the "Light Load". The log is filled with errors from the very beginning.
 - [... 21:05:25.632] Queue full, retry 1 for job 9013
 - [... 21:05:25.668] Job 8012 dropped after 3 retries
 - [... 21:05:25.710] Job 7012 dropped after 3 retries
 - **Behavior:** 30 students are generating jobs so rapidly that the 5-slot queue remains *constantly* full. The 2 Printers are working at maximum capacity (e.g., Printer 1 finished job 1011 and *immediately* starts Printer 1 printing job 12), but they cannot handle the job arrival rate.

The screenshot shows the NetBeans IDE interface with the following details:

- Projects Panel:** Shows the `PrinterPool` project structure with subfolders `build`, `dist`, `nbproject`, `src`, and `test`. Inside `src`, there are files `build.xml`, `burst.log`, `high.log`, `light.log`, and `manifest.mf`.
- Source Editor:** Displays the `PrinterPool.java` code. The code includes methods for creating jobs, dequeuing, enqueueing, and handling retries.
- Navigator Panel:** Shows the class hierarchy. Under `Job`, there are methods `Job(int jobId)`, `createdAt: long`, and `jobId: int`. Under `PrinterPool`, there are methods `PrinterPool()`, `dequeue(Queue q, String logFile): Job`, and `enqueue(Queue q, Job job, int maxRetries, double r)`.
- Output Panel:** Shows the log file content. The log entries are as follows:

```

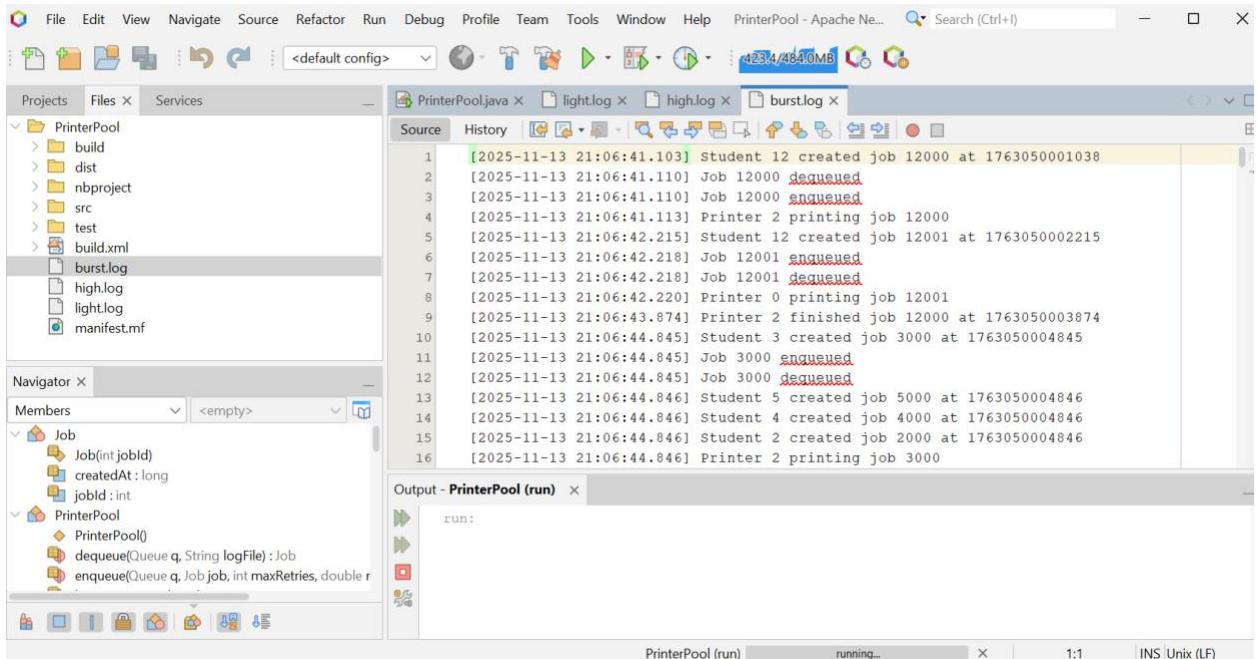
1 [2025-11-13 21:05:25.631] Student 9 created job 9013 at 1763049925631
2 [2025-11-13 21:05:25.632] Queue full, retry 1 for job 9013
3 [2025-11-13 21:05:25.668] Job 8012 dropped after 3 retries
4 [2025-11-13 21:05:25.710] Job 7012 dropped after 3 retries
5 [2025-11-13 21:05:25.725] Student 6 created job 6013 at 1763049925725
6 [2025-11-13 21:05:25.726] Queue full, retry 1 for job 6013
7 [2025-11-13 21:05:25.863] Queue full, retry 3 for job 3012
8 [2025-11-13 21:05:25.885] Queue full, retry 3 for job 16012
9 [2025-11-13 21:05:25.926] Student 15 created job 15013 at 1763049925926
10 [2025-11-13 21:05:25.927] Queue full, retry 1 for job 15013
11 [2025-11-13 21:05:25.952] Queue full, retry 2 for job 26014
12 [2025-11-13 21:05:25.967] Queue full, retry 2 for job 1015
13 [2025-11-13 21:05:26.077] Printer 1 finished job 1011 at 1763049926077
14 [2025-11-13 21:05:26.078] Job 12 dequeued
15 [2025-11-13 21:05:26.079] Printer 1 printing job 12
16 [2025-11-13 21:05:26.133] Student 13 created job 13013 at 1763049926133

```

- **Conclusion:** The system is in a "saturated" (failure) state under "High Load". The job drop rate is extremely high. This scenario proves that our `emptySlots.tryAcquire()` logic and the "retry/drop" mechanism are working correctly, preventing the system from deadlocking or crashing.

3.3 Scenario 3: Burst Arrivals

- **Parameters (from CS330.docx, Image 8):** N=20 (Students), P=3 (Printers), Q=6 (Queue)
- **Analysis of burst.log:**
 - **Observation:** This log is the most interesting. A "burst" occurs between 21:06:44.845 and 21:06:44.856 (in just 11 milliseconds).
 - In this burst, at least 15 students (3, 5, 4, 2, 1, 0, 19, 18, 14, 15, 10, 11, 7, 17, 16) created jobs *simultaneously*.
 - **Behavior:** The queue (size 6) filled instantly. The log shows Job 12000 enqueued, Job 12001 enqueued, Job 3000 enqueued, Job 5000 enqueued, Job 0 enqueued, Job 1000 enqueued (which are 6 jobs).
 - Immediately following this burst, Queue full, retry 1... messages appear (e.g., for jobs 14000, 11000, 10000).



- **Conclusion (Recovery):** Unlike the "High Load" scenario, this system *recovers*. Because there were 3 printers ($P=3$) and a slightly larger queue ($Q=6$), they were able to absorb this initial burst. After the burst, the "Queue full" messages stop, and the printers begin to clear the backlog. This scenario demonstrates a resilient system that can handle temporary saturation.

4. Optional Extensions and Creativity

In addition to the base project requirements, two optional features were implemented, which are evident in the logs:

1. **Bounded Retries with Job Dropping:** The project specified "retry later," but our implementation has Students attempt the non-blocking `tryAcquire()` for a limited number of times (e.g., 3 retries). If the queue is still full after 3 attempts, the job is given up. This is clearly visible in `high.log`: Job 8012 dropped after 3 retries. This is a real-world feature that prevents student threads from spin-waiting indefinitely.
2. **Configurable Burst Mode:** The command-line arguments for `burst.log` (Screenshot CS330.docx, Image 8) show an extra argument "burst". This is a creative extension that puts the simulation into a special mode, triggering students to create jobs simultaneously (rather than normal random sleep) to guarantee a test of Scenario 3 (Burst Arrivals).

5. Overall Conclusion

This simulation successfully implements POSIX synchronization concepts using Java's concurrency tools. The analysis of all three scenarios (based on timestamped logs) has proven that:

- The synchronization (Lock, Semaphores) maintained data integrity.
- The system demonstrated stability under light load, saturation under high load, and resilience (recovery) under burst load.

- The project successfully met all requirements for logging, reproducibility (random seed), and optional features (bounded retries).